



GFP: same 32 fitted training examples; exact replay through update 128; no dev/test

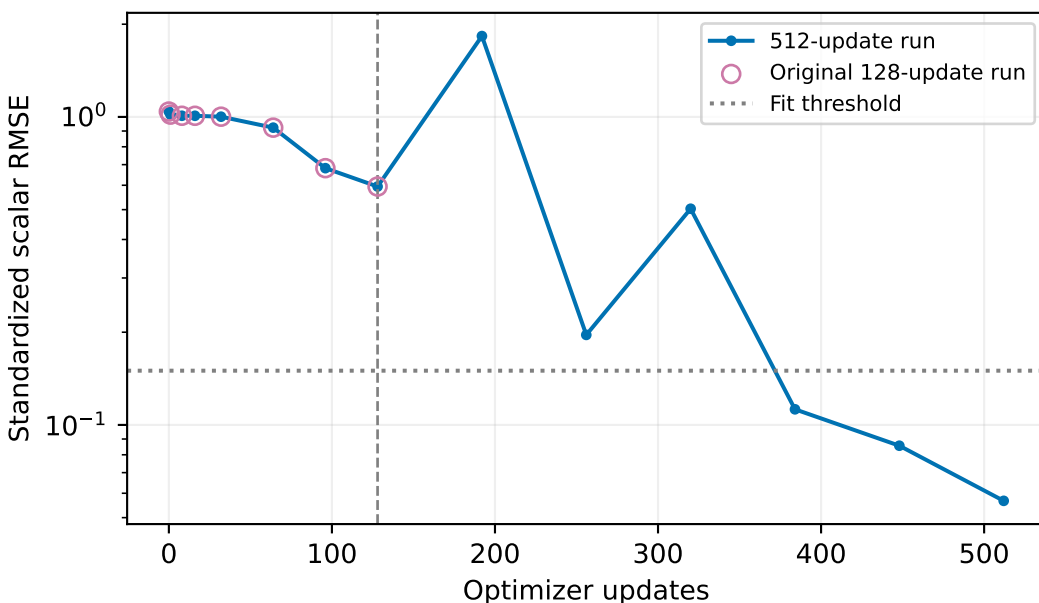
Five-bin accuracy



Categorical NLL



Standardized scalar RMSE



Optimization trace

